

Methodical Process

Dissection of PCAPS

- The first thing I wanted to do was view the packet fields
- I used tshark to view one of Ashley's PCAPS in XML format
 - `tshark -r pw1 -T pdml > pw1.xml`
 - PDML: Packet Description Markup Language
- I familiarized myself with the fields and tried to think of which would be useful
- Instead of picking random fields, I decided to research how to extract features from PCAP files
 - This is how I found Yang's PCAP study

Attempt at using netml tool

- Yang's study used a tool called netml to parse PCAP files and extract features
 - However, I was not able to get the netml tool to properly parse Ashley's PCAPS for an unknown reason
 - My only way forward was to replicate netml's functionality
- The netml tool has two parts, a PCAP parser and a machine learning tool
 - The machine learning functionality was not useful to me because I already had an Autoencoder to play with
 - I downloaded netml's parser code and began studying it

General concepts learned from netml parser code and Yang's study

- Group all packets together that are going from one machine to another
 - These are "flows" from source to destination
- Split these flows up into "subflows"
 - Fixed time intervals: Subflows are $\leq N$ seconds
 - Timeout period: Subflows are split using inter-arrival time $> N$ seconds
- Collect information about the packets in each subflow
 - This information becomes our feature vector
 - The study provides options for feature representation

Raw Field Extraction

- I was most comfortable with Yang's statistical representation of subflows
 - Time-related (subflow duration, packets/sec, bytes/sec)
 - Packet size stats (mean, std, min, max, q1, q2, q3)
- I revisited the PCAP XML file and picked fields necessary to obtain these statistics
 - This was done through inference
 - `frame.time_epoch`
 - This is the packet's capture time (or arrival time)
 - Needed for subflow durations
 - Time since Unix epoch (Jan 1, 1970)
 - A numpy function converts time epochs to seconds as necessary

- ip.len
 - Packet size in bytes
 - Specifically: Total length of IP frame
- For partitioning, Yang's study uses the following 5-tuple
 - (ip.src, srcport, ip.dst, dstport, ip.proto)
 - This means we must extract ports for both TCP and UDP
 - tcp.srcport, tcp.dstport
 - udp.srcport, udp.dstport
- Additionally, Yang's study notes the usefulness of TCP flags and IP TTL values
 - ip.ttl
 - tcp.flags
- I used tshark to extract the fields and save to CSV
 - tshark -r file.pcap -E header=y -E separator=, -E quote=d -E occurrence=f -T fields -e frame.time_epoch -e ip.len -e ip.proto -e ip.src -e ip.dst -e ip.ttl -e tcp.srcport -e tcp.dstport -e tcp.flags -e udp.srcport -e udp.dstport > file.csv

Combining Ashley's PCAPS

- Once I had the command for extraction, I needed to run it on all of Ashley's PCAPS
- I decided to combine all of the CSV files into a single CSV
 - I could have loaded all the CSV files up individually and concatenated them in pandas
 - However, I found it more straightforward to combine all the CSVs beforehand
- Preparation
 - All of Ashley's PCAPS were moved into a single folder
 - I renamed the files to make them easier to process with a shell script
 - pcap_one: pw1_1 . . . pw1_5
 - pcap_two: pw2_1
 - pcap_tree: pw3_1 . . . pw3_12
- Shell Script
 - After the renamed PCAPS were placed in a single folder, I created a shell script to run the tshark command on all of them
 - The script combines all CSV files for each group into one using Linux's cat command
 - cat pw1*.csv > combined_pw1.csv
 - cat pw3*.csv > combined_pw3.csv
 - The three CSV files are then concatenated into one large file
 - cat combined*.csv > combined_5G_pcaps.csv
 - All other csv files are removed

Data Preprocessing

- Yang's study is concerned only with IP packets, so I dropped all non-IP packets by checking for NaN values in the ip.proto column
- Remaining NaN values are converted to 0
 - ICMP packets do not have port numbers. A 0 should be a sufficient placeholder because it is part of the reserved port numbers and would not normally be used
 - Packets using UDP will have their TCP flags to 0, indicating no TCP flags

- The srcport and dstport columns are set to the corresponding TCP or UDP port numbers
 - This is done by taking the max of the TCP and UDP port
 - The original source and destination port columns (tcp.srcport, tcp.dstport, udp.srcport, udp.dstport) are dropped
- Since every CSV file has a header row, the concatenation of CSV files results in several header rows. All are dropped except the first one.
- All columns are then converted to an appropriate type
 - Strings for ID columns, datetime for time epochs, and integers for the rest
 - This is done because every column uses the raw “object” data type when first loaded

Flows

- Although Yang used the entire 5-tuple to partition the packets, I selected the tuple (ip.dst, ip.proto) for my experiments
 - I felt partitioning only by destination would better emulate “flows” into a device from multiple sources. This should help detect DDoS attacks
 - The destination port is removed so a multi-port attack can be detected, but this may not be necessary
- The flows are partitioned using Pandas functions
 - Unique Flow IDs are used to select matching rows (packets)
 - All packets matching the flow ID are set aside (partitioned)
 - A dictionary is used to link the flow IDs to their corresponding list of packets (partitions)
 - There must be at least 2 packets for flow analysis, otherwise the flow is dropped
 - See Supplemental.pdf for more info

Subflows

- Partitioning the data into long flows is not enough to accurately capture average network behavior
 - It also results in too few samples for machine learning
- The flows must be partitioned into smaller “subflows” using some method
- Yang’s study allows for subflow partitioning using either fixed time intervals ($\leq N$ seconds) or a timeout interval (inter-arrival time $> N$ seconds)
 - I created both methods for experimentation
 - I found that neither method resulted in significantly more samples, so I stuck with timeout interval because Yang used it for the statistical features portion of his study
- I iterated through the partitions (flows) and used Pandas functions to locate indices for subflow partitions
 - For timeout interval subflow selection, the primary functions are diff() and loc()
 - The supplemental PDF has more details on this process
- **Note that in a live system, there would be no partitioning into “subflows”**
 - We’d capture the last N seconds of traffic, partition it into flows, and extract statistical features
 - The autoencoder would tell us if the current network flow stats are normal or anomalous

Baseline Statistical Feature Selection

- As noted under the “raw field extraction” section, Yang’s study used 10 statistical features as a baseline
- I dropped subflow duration from the feature list
 - I don’t think network statistics need to be tied to specific durations
 - In a live system, there will only be a small window of network traffic analyzed at any time, perhaps just a few seconds
 - The stats during a 5 second capture may be normal, but if the autoencoder is trained to associate it with a 30 second flow, it will mark it as anomalous
 - What is important is that the stats are within our “normal” range, regardless of whether we analyze the last 5, 10, or 30 seconds in the real-world system

Additional Statistical Feature Selection

- Yang also noted that header information like TCP flags and IP time-to-live (TTL) can be useful for anomaly detection
 - Including information about ACK and SYN flags can help detect attacks like a SYN Flood
 - Time-to-live may be tied network topology, so differences here can indicate traffic from new devices.
 - Yang notes this may not be a useful feature if the TTL values are the same for all devices
- Yang appears to have tacked on the distribution of header values for each packet in a subflow to the end of his statistical feature vector
 - This results in a very long feature vector!
 - We’d also have to deal with uneven feature vectors because each subflow has a different number of packets
 - Yang discusses this in parts of the paper
- For ease, I chose the following statistical representations
 - ACK/sec
 - SYN/sec
 - Average TTL

Feature Extraction

- Pandas functions extract most of the features from the subflow partitions (which are Pandas dataframes)
 - The number of packets in a subflow is obtained using `shape[0]`
 - The duration is the difference from the last time epoch to the first
 - If there is only 1 packet in the subflow, it is skipped (same idea as before)
 - If the subflow duration is less than 1, it gets set equal to 1 for division purposes
- For ACK/sec and SYN/sec, I used bitmasks to isolate the individual flags before using Pandas `sum()` function to find the count
- Note that I changed the bytes/sec feature to MBits/sec because it is the standard representation of network traffic per second

Extractor Options

- As noted in the subflow section, my tool allows for partitioning by fixed interval or timeout period
 - Fixed interval is a better option when the flows do not have many breaks in them
 - Constant flows can't be broken up using a timeout period
- I also allow for the MBits/sec feature to be changed to KBits/sec
 - This was due to my Modbus experimentation
 - Network speeds were not high enough to justify using MBits
 - This is irrelevant to the 5G study
- Malicious subflows
 - These synthetic subflows are meant only for 5G testing
 - Based on ICMP ping flood
 - I used the minimum packet size (64 bytes)
 - The autoencoder does a good job of detecting these due to high packets per second
 - See the supplemental PDF for more info on malicious subflow generation

Autoencoder

- The autoencoder is a cross between Preston's autoencoder and Yang's autoencoder
- I copied Yang's default autoencoder structure
 - input layer d
 - hidden layer $d-1$
 - latent layer $[d/2]$
 - hidden layer $d-1$
 - output layer d
- Activations
 - Although Yang used LeakyReLU for his autoencoder, I ran some tests with Preston's original activations (ReLU, with Linear for output layer)
 - Preston's activations showed a better grouping of normal subflows than Yang's
 - I feel that we can more easily identify a clear threshold with these activations
- Optimizations
 - I attempted to implement some of the optimization ideas from the source I listed at the end of this paper
 - Tied Weights
 - Orthogonal weights
 - Unit norm constraint
 - No individual method or combination of these resulted in improved performance
 - I will have to rerun these experiments to provide specific details on the results of each
- Threshold
 - The remainder of the autoencoder reuses Preston's code
 - The threshold can be adjusted to improve the false positive rate, although it is already apparently low

Sources

<https://github.com/chicago-cdac/netml> (netml tool)

<https://arxiv.org/pdf/2006.16993> (Yang)

<https://arxiv.org/pdf/2004.13006>

- Unused but relevant
- Different “NetML”: Not actually the netml tool, but a dataset
 - They used an “Intel proprietary flow feature extraction tool”
- Discusses statistical flow features and other possible features
- Does not discuss Autoencoders explicitly, but mentions other machine learning algorithms including neural networks

<https://www.wireshark.org/docs/dfref/> (Wireshark reference page)

<https://towardsdatascience.com/build-the-right-autoencoder-tune-and-optimize-using-pca-principles-part-ii-24b9cca69bd6> (Autoencoder Optimization)